

Deep Learning Frameworks

使用不同的框架來實作訓練

- **Caffe**
- **Theano**
- **Torch**

Caffe

- **CIFAR-10**
- **60000 32X32 color images**
- **split into 10 classes**
 - airplane, automobile, bird, cat, deer
 - dog, frog, horse, ship, truck
- **`$caffe_path/data/get_cifar10.sh`**
- **`$caffe_path/examples/cifar10/create_cifar10.sh`**
- **Blob , Layer , Net , Solve**

Caffe - Blob -- 補充

The type of the data come from net is **Blob**:

```
{'prob': array([[ 0.,  0.,  0.,  0.,  0.,  
1.,  0.,  0.,  0.,  0.]], dtype=float32)}
```

We need to convert it to matrix that we can **use**:

```
[[ 0.  0.  0.  0.  0.  1.  0.  0.  0.  0.]]
```

This is second method that can extract any data from any **layer**:

```
[[ 0.  0.  0.  0.  0.  1.  0.  0.  0.  0.]]
```

Caffe - Layer -- 補充

```
layer {
  name: "conv1"
  type: "Convolution"
// type: Data, Convolution, Pooling, InnerProduct, ReLU, LRN,
//      DropOut, SoftmaxWithLoss, Accuracy
// type = Data: data (lmdb, leveldb), MemoryData
//      HDF5Data, ImagesData, WindowsData, DummyData

  bottom: "data" # 輸入層：數據層
  top: "conv1" # 輸出層：卷積層1

  # 濾波器 (filters) 的學習速率因子和衰減因子
  param { lr_mult: 1 decay_mult: 1 }

  # 偏置項 (biases) 的學習速率因子和衰減因子
  param { lr_mult: 2 decay_mult: 0 }

  convolution_param {
    num_output: 96 # 96個濾波器 (filters)
    kernel_size: 11 # 每個濾波器 (filters) 大小為11*11
    stride: 4 # 每次濾波間隔為4個像素
    weight_filler {
      type: "gaussian" # 初始化高斯濾波器 (Gaussian)
      std: 0.01 # 標準差為0.01，均值默認為0
    }
    bias_filler {
      type: "constant" # 初始化偏置項 (bias) 為零
      value: 0
    }
  }
}
//source code: src/layers/convolution_layer.cpp | cu
```

Caffe - net

```
...
layer {
  name: "cifar"
  type: "Data"
  top: "data"
  top: "label"
  transform_param {
    mean_file: "examples/cifar10/mean.binaryproto"
  }
  data_param {
    source: "examples/cifar10/cifar10_train_lmdb"
    batch_size: 100
    backend: LMDB
  }
}
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  ...
}
...

```

Caffe - solver

```
# reduce the learning rate after 8 epochs (4000 iters) by a factor of 10

# The train/test net protocol buffer definition
net: "examples/cifar10/cifar10_quick_train_test.prototxt"
# test_iter specifies how many forward passes the test should carry out.
# In the case of MNIST, we have test batch size 100 and 100 test
iterations,
# covering the full 10,000 testing images.
test_iter: 100
# Carry out testing every 500 training iterations.
test_interval: 500
# The base learning rate, momentum and the weight decay of the network.
base_lr: 0.001
momentum: 0.9
weight_decay: 0.004
# The learning rate policy
lr_policy: "fixed"
# Display every 100 iterations
display: 100
# The maximum number of iterations
max_iter: 4000
# snapshot intermediate results
snapshot: 4000
snapshot_prefix: "examples/cifar10/cifar10_quick"
# solver mode: CPU or GPU
solver_mode: GPU
```

Caffe - train

```
cd $caffe_path  
./build/tools/caffe train  
--solver=./examples/cifar10/cifar10_quick_solver.prototxt
```

```
#resume training from the half-way  
point snapshot  
caffe train -solver  
examples/mnist/lenet_solver.prototxt  
-snapshot  
examples/mnist/lenet_iter_5000.solverstate
```

Caffe - time -- 補充

```
caffe time -model examples/mnist/lenet_train_test.prototxt
-iterations 10
```

```
I1005 05:58:47.061524 1233 caffe.cpp:374] *** Benchmark begins ***
I1005 05:58:47.061544 1233 caffe.cpp:375] Testing for 10 iterations.
I1005 05:58:47.115362 1233 caffe.cpp:403] Iteration: 1 forward-backward time: 53 ms.
I1005 05:58:47.169019 1233 caffe.cpp:403] Iteration: 2 forward-backward time: 53 ms.
I1005 05:58:47.222576 1233 caffe.cpp:403] Iteration: 3 forward-backward time: 53 ms.
I1005 05:58:47.275920 1233 caffe.cpp:403] Iteration: 4 forward-backward time: 53 ms.
I1005 05:58:47.329260 1233 caffe.cpp:403] Iteration: 5 forward-backward time: 53 ms.
I1005 05:58:47.383637 1233 caffe.cpp:403] Iteration: 6 forward-backward time: 54 ms.
I1005 05:58:47.438005 1233 caffe.cpp:403] Iteration: 7 forward-backward time: 54 ms.
I1005 05:58:47.491442 1233 caffe.cpp:403] Iteration: 8 forward-backward time: 53 ms.
I1005 05:58:47.544661 1233 caffe.cpp:403] Iteration: 9 forward-backward time: 53 ms.
I1005 05:58:47.598136 1233 caffe.cpp:403] Iteration: 10 forward-backward time: 53 ms.
I1005 05:58:47.598157 1233 caffe.cpp:406] Average time per layer:
I1005 05:58:47.598170 1233 caffe.cpp:409]     mnist forward: 0.0107 ms.
I1005 05:58:47.598176 1233 caffe.cpp:412]     mnist backward: 0.0006 ms.
I1005 05:58:47.598179 1233 caffe.cpp:409]     conv1 forward: 5.9642 ms.
I1005 05:58:47.598182 1233 caffe.cpp:412]     conv1 backward: 5.9309 ms.
I1005 05:58:47.598186 1233 caffe.cpp:409]     pool1 forward: 2.7406 ms.
I1005 05:58:47.598188 1233 caffe.cpp:412]     pool1 backward: 0.5376 ms.
I1005 05:58:47.598191 1233 caffe.cpp:409]     conv2 forward: 10.0786 ms.
I1005 05:58:47.598196 1233 caffe.cpp:412]     conv2 backward: 19.9588 ms.
I1005 05:58:47.598208 1233 caffe.cpp:409]     pool2 forward: 1.6229 ms.
I1005 05:58:47.598212 1233 caffe.cpp:412]     pool2 backward: 0.686 ms.
I1005 05:58:47.598214 1233 caffe.cpp:409]     ip1 forward: 1.9727 ms.
I1005 05:58:47.598217 1233 caffe.cpp:412]     ip1 backward: 3.7615 ms.
I1005 05:58:47.598220 1233 caffe.cpp:409]     relu1 forward: 0.0241 ms.
I1005 05:58:47.598224 1233 caffe.cpp:412]     relu1 backward: 0.0282 ms.
I1005 05:58:47.598227 1233 caffe.cpp:409]     ip2 forward: 0.1135 ms.
I1005 05:58:47.598230 1233 caffe.cpp:412]     ip2 backward: 0.1435 ms.
I1005 05:58:47.598239 1233 caffe.cpp:409]     loss forward: 0.0371 ms.
I1005 05:58:47.598243 1233 caffe.cpp:412]     loss backward: 0.0018 ms.
I1005 05:58:47.598250 1233 caffe.cpp:417] Average Forward pass: 22.5693 ms.
I1005 05:58:47.598254 1233 caffe.cpp:419] Average Backward pass: 31.0536 ms.
I1005 05:58:47.598258 1233 caffe.cpp:421] Average Forward-Backward: 53.6 ms.
I1005 05:58:47.598263 1233 caffe.cpp:423] Total Time: 536 ms.
I1005 05:58:47.598264 1233 caffe.cpp:424] *** Benchmark ends ***
```


Caffe - Python - API

```
# Import required Python libraries
%matplotlib inline
import os
import numpy as np
import matplotlib.pyplot as plt
import caffe
import random

# Choose network definition file and pretrained network binary
MODEL_FILE = '/home/ubuntu/caffe/examples/cifar10/cifar10_quick.prototxt'
PRETRAINED = '/home/ubuntu/caffe/examples/cifar10/cifar10_quick_iter_4000.caffemodel'

# Load a random image
x = caffe.io.load_image('/home/ubuntu/caffe/examples/images/' + str(random.randint(1,18)) + '.png')

# Display the chosen image
plt.imshow(x)
plt.axis('off')
plt.show()

# Load the pretrained model and select to use the GPU for computation
caffe.set_mode_gpu()
net = caffe.Classifier(MODEL_FILE, PRETRAINED,

mean=np.load('/home/ubuntu/caffe/caffe/examples/cifar10/cifar10_mean.npy').mean(1).mean(1),
               raw_scale=255,
               image_dims=(32, 32))

# Run the image through the pretrained network
prediction = net.predict([x])

# List of class labels
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

# Display the predicted probability for each class
plt.plot(prediction[0])
plt.xticks(range(0,10), classes, rotation=45)
# Display the most probable class
print classes[prediction[0].argmax()]
```

Theano

- **MNIST**
- **28X28 black and white images**
- **handwritten digits**
- <http://yann.lecun.com/exdb/mnist/>

Theano - Python - 補充

- <http://deeplearning.net/tutorial/lenet.html#lenet>
- https://github.com/lisa-lab/DeepLearningTutorials/blob/master/code/convolutional_mlp.py
- **nvidia-docker**
- **nvidia-docker run -it nakosung/dockerfiles-1**
- **nvidia-docker run -it nvidia/cuda /bin/bash**

```
wget https://repo.continuum.io/archive/Anaconda2-5.0.0.1-Linux-x86_64.sh
bash Anaconda2-5.0.0.1-Linux-x86_64.sh
conda install theano
git clone https://github.com/lisa-lab/DeepLearningTutorials.git
./DeepLearningTutorials/data/download.sh
cp DeepLearningTutorials/data/mnist-* DeepLearningTutorials/code/
cd DeepLearningTutorials/code/
python mlp.py
Using gpu device 0: GeForce GTX 1060 6GB (CNMeM is enabled with initial size: 10.0% of
memory, CuDNN 5105)
/usr/local/lib/python2.7/dist-packages/theano/sandbox/cuda/__init__.py:600:
UserWarning: Your CuDNN version is more recent than Theano. If you see problems, try
updating Theano or downgrading CuDNN to version 4.
warnings.warn(warn)
... loading data
... building the model
... training
epoch 1, minibatch 2500/2500, validation error 9.620000 %
    epoch 1, minibatch 2500/2500, test error of best model 10.090000 %
epoch 2, minibatch 2500/2500, validation error 8.610000 %
    epoch 2, minibatch 2500/2500, test error of best model 8.740000 %
epoch 3, minibatch 2500/2500, validation error 8.000000 %
    epoch 3, minibatch 2500/2500, test error of best model 8.160000 %
```

Torch

- **char-rnn** 程式與實作
<https://github.com/karpathy/char-rnn>
- 資料集
<http://cs.stanford.edu/people/karpathy/char-rnn/>
- 研究分析
<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Torch 實作

- **nvidia-docker**
- **nvidia-docker run -it kaixhin/cuda-torch:8.0**
cd
git clone <https://github.com/karpathy/char-rnn.git>
th train.lua
th sample.lua cv/lm_lstm_epoch9.46_1.4349.t7
th convert_gpu_cpu_checkpoint.lua
cv/lm_lstm_epoch30.00_1.3950.t7
- **PyTorch <https://github.com/nearai/pytorch-tools>**
- **torch lenet**
<https://github.com/NVIDIA/DIGITS/blob/master/digits/standard-networks/torch/lenet.lua>